

BTC15 V2

Offline correction candidate

Release review | 04 October 2026 | Product R1

NOT DEPLOYED. NOT APPROVED FOR PRODUCTION.

SIGNAL ONLY / MANUAL EXECUTION ONLY / NO ORDERS

The coordinated candidate implements the approved product/evidence revision and offline source-bootstrap preparation. Frozen strategy files, model inputs/weights, thresholds, freshness, eligibility and protection/EXIT policy remain unchanged. The existing scoring bridge remains the only scorer.

Evidence	Result
MAIN offline qualification	280 tests passed
V8.1 current source contract	19 tests passed on the V8.1 candidate
Existing bridge and descriptive diagnostics	17 tests passed
Full assembled dashboard DOM	12 scenarios passed; four scripts executed; zero legacy action-card writes
Strict end-to-end fixture workflow	1 then 2 fully scoreable settled fixtures; production claim remains blocked

Open gates are explicit

Rendered browser preview was blocked by browser security policy. Physical iPhone/iPad Safari behavior is not accepted. D04 production stalls were not reproduced offline, so no performance optimization was selected. Real quote-lag attribution, one real strictly scoreable settled contract, later accumulation and user alert delivery remain live acceptance work.

No production deployment, restart, task edit, service retirement, order or user notification occurred. Both commits remain local. The release archive contains exact source additions, incremental git bundles, patches, full logs, original defect matrix, candidate dispositions, inactive alert drafts, rollback and acceptance instructions.

Read this report with BTC15_V2_PRODUCT_RELEASE_REVIEW.md and RELEASE_INDEX.json in the archive. The full 18-defect matrix retains every original evidence/dependency/semantic-impact field; pages 4-9 summarize the candidate disposition.

Frozen semantics and approved boundaries

Byte identity and runtime equivalence

All 697 pre-existing MAIN tracked files and 607 pre-existing V8.1 tracked files match their frozen commits. Original lane-specific freeze verification stays mandatory. Additional source/model/dependency hashes protect integration assumptions. Differential replay checks frozen MAIN native functions/state and V8.1 iterations, original offer counts, sleep cadence and complete processor record/state outputs.

Policy	Preserved
EARLY	Ask ≤ 45 cents; fair $\geq 75\%$; 2-10 minutes; absolute gap $\geq \$25$; original confirmations and protections
FINAL	Independent qualification, thresholds and confirmation/protection behavior
SCALP	Actual original ask $\rightarrow$ strictly later same-side bid; +5-cent arm, 4-cent giveback, 180-second policy
Source freshness	BRTI $\leq 5s$; BTC $\leq 10s$; quote $\leq 6s$; capture $\leq 15s$; original close-time and causal bounds
Performance classification	Original sticky unavailable exclusion and both-lane strict conjunction retained

D02: source preparation only

At most two instances of the original timestamped provider; next-contract preparation begins no earlier than 30 seconds before open. This is a transport window, not a strategy gate. The current live book is never reset for preparation. Selection re-fetches the exact official ticker/target/window; staged target is never used. No pre-open qualification, REST/last-price fallback, extra frame or faster authoritative evaluation is allowed. Tests cover delayed target, identity drift, bootstrap, source causality and bounded rollover.

D13: separate evidence revision

New data belongs to a distinct revision/deployment root on the existing volume. Build/deployment metadata, runtime epochs and excluded startup/restart slots prevent silent mixing. Old evidence remains untouched. Administrative attempts never enter Directional.process or Scalp.process and never count as scored observations. A reviewed control-plane receipt is required for the new production cohort; fixtures and runtime self-reports are insufficient.

Evidence: main_semantic_inventory.json; v81_semantic_inventory.json; offline_qualification_tests.log; v81_current_source_tests.log; BTC15_V2_PRODUCT_RELEASE_R1.json.

Measurement and qualification limits

D04 was measured before selecting optimization

Historical evidence remains 49 fresh-at-cutoff decisions followed by delayed publication, 9.322-24.470 seconds, median 15.168. Five other internal gaps overlap quote waits. These measurements do not identify the inner operation.

With exact pinned dependencies, 50,416 candle rows and 241 native ticks, the first offline stage profile ran 40 samples. Fair-frame median was 17.959 ms; frozen RF prediction median 26.818 ms. That probe excluded the optional legacy shadow model. A follow-up reconstructed its original startup training from the exact 1,474-row frozen CSV: ten complete loops with the model took 48.226-59.950 ms. One standalone prediction sample reached 74.684 ms. No multi-second production stall was reproduced.

No model, thread, fsync, scheduler or strategy-clock optimization is included. Stage timing is administrative instrumentation, not proof that the production bottleneck is repaired. Local CPU, network and disk conditions differ from production.

Assembled dashboard evidence

All four assembled scripts ran in JSDOM, including legacy timers/errors. Tests verify sole action-card ownership, PASS prominence despite 95% model probability, real frozen later-bid EXIT output, independent SCALP freshness, malformed/identity rejection, cached timestamp expiry, abort fallback, old-reply rejection and hide across rollover. A comparable synthetic 3-cent accepted-book move reaches the displayed prices with zero post-render delta. The six-second lease stays intact.

JSDOM has no layout engine and is not Safari. Local HTTP and file previews were both blocked by the browser security policy; no workaround was attempted. Responsive rendered layout and physical Safari acceptance are unverified. Production 2-3-cent lag has not been exactly attributed.

Legacy test debt is not hidden

The older V8.1 module has 18 tests: 14 pass, four error. The identical four errors reproduce at the exact frozen baseline: obsolete publish-function fixture calls and a missing required BTC source timestamp. The original file is preserved. Nineteen current API/source-contract tests pass without weakening runtime gates. Logs for both old runs and the current suite are included.

Fixture success is functional evidence only

Original processors, durable journals, consistent snapshots and the established bridge produced 1 then 2 strictly scoreable settled fixture contracts. The bridge still returns SYNTHETIC_FIXTURE_ONLY and blocks a production claim. No actual accepted entry, fill or performance reliability is inferred.

18-defect disposition | 01-03

D01 Skipped MAIN attempts leave no V2 observation.

IMPLEMENTED OFFLINE

Root cause: The ladder hook is appended only to the successful native try-body. NO ACTIVE and QUOTE WAIT continue; outer exceptions also bypass it. SQLite can be intact while runtime evidence is incomplete.

Candidate evidence: Skipped native attempts, exceptions, source clocks and accepted offers have separate bounded evidence. Actual worker test proves one strategy commit and three linked administrative rows.

Remaining gate: Hard-kill loss of an uncommitted administrative item remains possible; no claim of lossless administrative capture.

D02 Openings start late; every new V8.1 ticker incurs an unavailable observation.

IMPLEMENTED OFFLINE CONDITIONAL SOURCE PREPARATION

Root cause: MAIN performs unopened discovery and open-list I/O before staged exact fetch; target may not yet exist. Both quote providers reset on a new ticker and return None. V8.1 immediately journals that forced warm-up as UNAVAILABLE. Canonical V8.1 discovery can also return no active market.

Candidate evidence: Same decoder/provider with exact official handoff, target/window invariants, bounded two-instance pool, no pre-open or extra native evaluation; delayed target/reset/reconnect/causality tests.

Remaining gate: Live readiness is not guaranteed. A missing official target or post-open source still excludes strict scoreability.

D03 Closing-edge observations disappear or record unavailable books.

REPRESENTED NOT CURED

Root cause: A paired executable book can genuinely become unavailable near close. MAIN skips publication on this condition; V8.1 records it. Missing executable sides are not packet-loss evidence.

Candidate evidence: Empty paired book removes prices and actions; native attempt is accounted without synthetic strategy frames. Original empty/sequence/stale gates retained.

Remaining gate: External executable-book unavailability cannot be corrected by this implementation.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

18-defect disposition | 04-06

D04 49 historical internal gaps are slow post-input processing, not logged waits.

MEASURED NO OPTIMIZATION

Root cause: A native decision has already acquired fresh BRTI, but completion and ladder publication occur much later. The critical path includes fair-frame/model work, native shadow lifecycle work, information anchor copying and synchronous cohort fsync before the ladder hook. Existing logs do not time these individually.

Candidate evidence: 40-stage-sample benchmark and 10 complete loops with original legacy shadow model; 48.226-59.950 ms locally. Stage evidence added.

Remaining gate: Production 9.322-24.470-second delays were not reproduced. Inner cause unresolved; no scheduler/model/fsync/thread optimization selected.

D05 Fresh action snapshots expire between MAIN publications.

IMPLEMENTED PRESENTATION TRUE EXPIRY REMAINS

Root cause: MAIN nominal evaluation interval is 5s while BRTI is already aged when sampled. The action expires at $\min(\text{close}, \text{BRTI}+5, \text{BTC}+10, \text{quote}+6, \text{capture}+15)$. Faster GETs cannot renew that evidence. Network time further reduces usable display life.

Candidate evidence: Typed expiry withdraws guidance; stable DOM structure and changed-value writes; source budgets unchanged. Cached responses cannot renew leases.

Remaining gate: Native five-second clock can still outrun remaining BRTI lifetime. Continuous actionable display is not promised.

D06 Competing legacy and V2 writers flash the same action cards.

IMPLEMENTED OFFLINE

Root cause: Installer removes the inline V8.1 script but retains legacy applyState, markUnavailable and the one-second countdown. They write FINAL/EARLY/SCALP IDs also written every 250ms by V2. Calling V2 at the end of successful applyState does not retire legacy ownership.

Candidate evidence: All legacy scripts execute; zero legacy writes to owned action fields across success, failure, timers and lifecycle in the assembled DOM test.

Remaining gate: Rendered browser and physical Safari validation remain pending.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

18-defect disposition | 07-09

D07 Specific unavailable reasons and historical origins are discarded by the panel.

IMPLEMENTED OFFLINE

Root cause: SOURCE_EXPIRED responses omit candidate, served_ts and expires_at. Recorded rejected frames often have no expiry. The panel requires these fields for all responses, throws UNQUALIFIED, clears cache and renders generic waiting text. The server intended to retain historical origin on expiry.

Candidate evidence: Typed AVAILABLE/UNAVAILABLE validation; safe reason and identity-only historical context retained; malformed payload fails closed.

Remaining gate: Historical origin is never current entry permission.

D08 Fresh SCALP guidance is hidden when MAIN expires.

IMPLEMENTED OFFLINE

Root cause: aligned requires qualified MAIN before it will show a qualified V8.1 record. MAIN action freshness is being used as a prerequisite for another independently qualified lane. It is not a SCALP strategy gate.

Candidate evidence: SCALP validates its own source deadline against exact current official identity even when MAIN action expires. Wrong target and rollover rejected.

Remaining gate: No SCALP strategy gate changed; requires verified identity and actual lane freshness.

D09 Probability dominates PASS; stale direction text and obsolete production label persist.

IMPLEMENTED OFFLINE

Root cause: Large side/confidence typography is reused for unavailable output. Unavailable rendering omits titles/arrows and context cleanup. Legacy sets DIAGNOSTIC V13.2-P1. The raw model estimate is visually stronger than FINAL qualification.

Candidate evidence: FINAL PASS/UNLOCKED is dominant, raw model probability secondary, stale labels/arrows cleared, obsolete production diagnostic label removed.

Remaining gate: Visual typography and mobile overflow require rendered inspection.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

18-defect disposition | 10-12

D10 Mobile hide/resume and Safari behavior are not accepted end-to-end.

DOM QUALIFIED BROWSER BLOCKED

Root cause: Both page scripts own lifecycle invalidation. V2 uses busy flags, generation tokens and AbortSignal.timeout; pending requests across suspension need explicit cancellation/reset. Existing test covers only the isolated panel in a mock DOM, not the assembled page or Safari.

Candidate evidence: AbortController timeout fallback, generation tokens, hide/resume, bfcache, cached expiry and rollover tests pass with full assembled scripts.

Remaining gate: Browser security policy blocked local HTTP and file preview. JSDOM is not a layout engine or Safari; physical iPhone/iPad pending.

D11 Reported 2-3 cent quote lag has no exact DOM-to-book measurement.

IMPLEMENTED MEASUREMENT LIVE ATTRIBUTION PENDING

Root cause: Displayed UP/DOWN asks use the native snapshot and approximately 5s legacy refresh; their DOM lacks the accepted quote timestamp/sequence. The information route already exposes timestamped same-provider quotes, but the browser comparison is not paired atomically to the latest accepted book.

Candidate evidence: Same accepted tuple projects four prices and source clocks; 500-ms presentation polling; bounded DOM receipts. Injected comparable 3-cent move reaches DOM with zero post-render delta.

Remaining gate: Exact live 2-3-cent lag attribution and display latency require timestamped comparable production tuples; no invented production measurement.

D12 Strict scoreability is permanently vetoed by any recorded unavailable state.

STRICT PRESERVED DIAGNOSTICS ADDED

Root cause: Journal missing is sticky after any unavailable_reason. Bridge independently excludes unavailable intervals and ANDs both lanes. V8.1 forced ticker bootstrap (D02) therefore excludes every newly observed complete contract under the frozen implementation. Administrative attempt records alone cannot create a strict scoreable contract.

Candidate evidence: Original strict functions are AST-identical; established scorer byte-identical. Descriptive recording/source axes remain separate. Fixture 1 then 2 strict contracts uses original processors and bridge.

Remaining gate: Any natural unavailable record still excludes strict cohort. Fixtures never count as live performance.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

18-defect disposition | 13-15

D13 A corrected deployment cannot simply reuse the frozen scorer identity or mixed journal.

APPROVED REVISION IMPLEMENTED OFFLINE

Root cause: Runtime startup verifies frozen file hashes. Bridge allowlists exact builds/deployments/volumes and rejects old-build event rows. Replacing build constants while retaining old events in the same volume would fail or tempt silent cohort mixing. Process restart also resets continuity and can preserve historical origin.

Candidate evidence: Exact manifest/build/service/volume binding, new deployment root, immutable metadata and runtime epochs, excluded startup/restart slots, foreign root byte-preservation and mixed-build rejection tests.

Remaining gate: Actual production receipt cannot exist before approved deployment; runtime self-report does not authorize production scoring.

D14 Pending settlement and absent-lane tails are obscured by partial status.

IMPLEMENTED OFFLINE

Root cause: Partial takes precedence over pending in the headline. Snapshots were taken at different times. Settlement reader correctly requires finalized binary official results and retries unresolved receipts after 300s on a 30s cycle.

Candidate evidence: Settlement state and absent-lane snapshot tails reported independently of partial source/observation state; original official finalized-only linkage preserved.

Remaining gate: No receipt or finalization speed is fabricated; pending remains pending.

D15 Existing alert tasks use an older first-read handoff; delivery is unproven.

DRAFT ONLY PENDING POST CORRECTION

Root cause: Both enabled tasks point first to ops/btc15-v2-qualification-20261003, while the verified live scoring handoff is ops/v2-score at 61b3b4d. Task enablement and last execution do not prove user push/email receipt.

Candidate evidence: Draft updates preserve the two existing task IDs, schedules, rules and suppression. Prior prompts retained as evidence. No task mutation/delivery performed.

Remaining gate: After approval and actual correction, verify settings/destination and obtain Eric receipt from one labeled existing-path test.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

18-defect disposition | 16-18

D16 Broad service inventory and error noise are unresolved dependencies, not cleanup permission.

INVENTORY ONLY

Root cause: The environment contains many research/legacy processes. Native decision-path work and high dashboard read traffic coexist. Resource summaries and broken pipes alone do not identify a guilty service or justify retirement.

Candidate evidence: No services retired, no paid resources created, no cost/savings assertion. Native client disconnect handled as bounded handler return.

Remaining gate: Resource contention and dependency inventory do not establish a guilty service.

D17 Retention and observation/commit health need distinct accounting.

IMPLEMENTED BOUNDS WORKFLOW PENDING LIVE

Root cause: Retention is 30 days OR 300,000 events, whichever is reached first. At ~1 V81 event/s the count cap is about 3.47 days before admin/settlement overhead. HTTP coverage exposes only 96 contracts. Commit counters do not count skipped native attempts.

Candidate evidence: Scored retention unchanged; retained sequences in existing snapshots; separate bounded admin health/range/bytes. Read-only consistent snapshot and same workflow in fixtures.

Remaining gate: 30 days OR 300000 scored events; approximately 3.47 days at 1 Hz. Admin cap can be earlier. 96-contract coverage is not full history; live retention check pending.

D18 Source rejection labels are too coarse to assign every failure to the external source.

IMPLEMENTED ADMINISTRATIVE CLOCKS

Root cause: BRTI_SOURCE_UNAVAILABLE and BTC_SOURCE_UNAVAILABLE combine readiness/causal-age conditions. Some source-valid decisions age during processing (D04). Quote availability also conflates empty side, stale time, sequence state and bootstrap. No evidence supports a general packet-loss claim.

Candidate evidence: Separate source/received/cut/capture/processing/commit/serve clocks, known safe rejection subreasons and expired-during-processing markers; missing clocks stay null.

Remaining gate: No receipt-time substitution and no packet-loss inference; some upstream receipt clocks are not available in frozen data.

Exact historical evidence, dependencies, severity, semantic impact, tests and deployment requirements remain in qualification/defect_matrix_with_candidate_disposition.json and prior_audit/defect_matrix.json.

Ordered release and rollback

1. Review and close offline gates

Review both commits, manifest and unchanged semantic inventory. Finish rendered assembled-page and iPhone/iPad Safari checks in an authorized environment. Keep D04 attribution unresolved until measured. No implicit gate waiver.

2. Obtain explicit production approval

Approval must name these exact lane commits and manifest. Current approval permits offline implementation/testing only. Any additional source/evidence/strategy-definition change requires review.

3. Preserve identities and evidence

Capture existing start commands/configuration and verify the existing volumes; keep secrets out of the report. Take consistent read-only old journal snapshots. Inspect auto-deploy behavior before any authorized source push.

4. Apply one coordinated release after approval

Existing two production services only. Original volume guards plus the reviewed launcher. Pin approved manifest/build per lane; preserve MAIN information-export opt-in. No shared BRTI-owner change, new service, cleanup or order path.

5. Bind actual deployments

Record actual control-plane deployment IDs, commits, service/project/environment/volumes and runtime starts. Fill and hash the receipt template only then. Startup pinning is not production scoring authorization. New cohort starts at the first full window strictly after both starts.

6. Accept, settle, score and verify delivery

Use one bounded live window plus rollover, then fresh detached paired snapshots and the unchanged strict bridge. Confirm a later checkpoint under the same receipt. Update the same two alert tasks only after the verified handoff; require actual Eric receipt.

Rollback is concrete

Restore exact frozen MAIN de4f3e20... and V8.1 60e6ebdd... commits and original start commands given verbatim in the release review. Restore only changed settings from captured values. Preserve both old/new roots and snapshots; never merge them. Rollback creates new deployment IDs, so historical allowlists must not silently admit rollback events. Stop new-cohort scoring, restore prior alert prompts only if activated, and re-verify source identities/frozen hashes/no orders.

Bounded live acceptance after approval

Official identity and causal inputs

Exact official ticker, fixed target and open/close in both lanes. BRTI/BTC/quote source times causal at native cuts and publication. No pre-open frame, faster strategy evaluation or synthetic observation.

Readable action cards and real expiry

Observe EARLY/FINAL/SCALP with PASS dominant where unqualified. Source expiry must remove actionable guidance; fresh SCALP may remain when its own evidence and exact identity are current. PASS/no entry is legitimate: do not manufacture signals.

Comparable quote timeliness

Pair latest accepted book with DOM using identical ticker/epoch/sequence. Record four prices, exchange/accepted/published/served/render clocks, source age, latency bounds and price delta. A new 2-3-cent move should appear at the next successful display poll/render without renewing the original six-second lease. Do not compare mismatched tuples or infer one-way network latency from unsynchronized clocks.

iPhone/iPad Safari and rollover

Hide/resume during the window and across rollover. Abort old requests; reject previous-generation replies; require fresh official identity on return. Check actual layout, no stale direction/origin and readable unavailable state.

Persistent evidence and D04 timing

Reconcile observed START/FINISH attempts, native offers, strategy commits and served views. Measure native attempt/publication cadence and cutoff-to-commit distributions with worst gaps. Administrative health is separate, bounded and non-authoritative. Do not infer packet loss.

Strict settled result and later accumulation

Wait for finalized official settlement. Snapshot each lane via SQLite online backup, transfer both .sqlite3 and .sqlite3.receipt.json, run the existing bridge with exact reviewed release receipt/hash. Require at least one real fully scoreable settled contract. If blocked, report the reason and stop; no relaxed gates or new multi-day campaign. Confirm a later bounded checkpoint under the same workflow.

Alert delivery and retention

Use the existing two tasks and existing delivery path; verify settings/destination and obtain Eric receipt of one labeled test. No duplicate watchers. Scored retention remains 30 days OR 300,000 events (~3.47 days at 1 Hz); 96-contract coverage is not history. Admin has separate 7-day/100,000-record/64-MiB bounds. Cleanup remains inventory-only.

Exact release identity

MAIN candidate commit

020f8a38d07e0a11e983a1e3bee3b06a4898c2f3

Frozen parent: de4f3e20b8657eb8cfee91bd4e525c103b5bf513

Incremental bundle SHA256: 3e71a3e9b99bcda2187bcd2eeb702b08d2615c2191e13b6e8fbd815c153097c9

V81 candidate commit

07406e46945cc0bfefbbb25299c5da53e87ac6a2a

Frozen parent: 60e6ebdd03e81a4c84385e1b5122302f3cf0b9f1

Incremental bundle SHA256: d1aad2a367e818cde37e28fffc8fe9092b0f78bb6e1301ab83f53f3ea163cac8

Content manifest SHA256

7289357977b4771b16644b32f75101883a4ecaaf7e8f82eef8d8042848d040f0

Established scorer SHA256

9812c23a6200bf31c2c5eb4e8f317fe32cfcafeaa37d2a9b9a041967447d1a

Original bridge / live handoff

1d942014af22ab4acb3048a29dbbc831d5ecc924
61b3b4ddb3a816559355802ddd8dba27abfcdf00

Archive contents

RELEASE_INDEX.json and SHA256SUMS; per-lane incremental bundles and binary patches; exact committed candidate sources; release manifest and review instructions; invalid deployment-receipt template; original full audit and repair plan; 18-defect candidate matrix; all qualification logs/profiles; offline assembled preview; inactive draft updates to the same two alert IDs.

The receipt template has no actual deployment IDs or starts and cannot authorize production. Bundle verification requires the frozen parent commit. Use the provided restore/review instructions in an isolated checkout. No GitHub push or production deployment has occurred.

Decision requested next

Review this exact candidate and the remaining gates. Production remains stopped at the approval boundary. This package is not a claim that the live product correction or final acceptance is complete.